14 Administration

This chapter covers

- Horizontally scaling the cluster
- Internode communication
- Shards and replica sizing
- Working with snapshots and restoration
- Advanced configuration
- Understanding the master role in a cluster

So far, we've seen the inner workings of Elasticsearch, including its excellent queries and other features. We haven't bothered with advanced configurations like how nodes communicate with each other, how big shards need to be, and what settings to change to modify the Kibana port. In this chapter, we discuss these administrative features. We address some of them by using the queries we created while running the search server.

One powerful feature of Elasticsearch is its ability to scale up the server to provide petabytes of data. There is no complexity in setting this up other than procuring additional nodes. We cover how to scale the cluster in the first part of this