Go to the second tab, Repositories, and click Register a Repository. The next page is shown in figure 14.11.

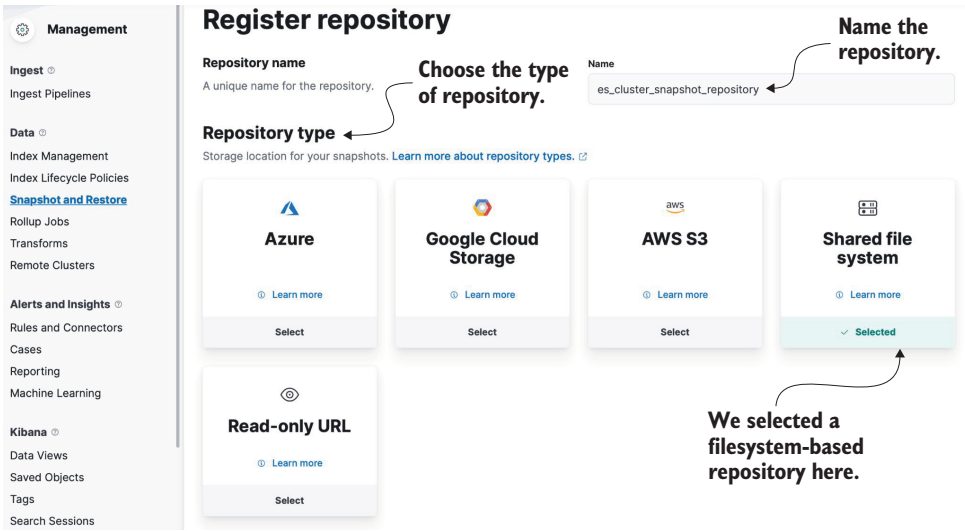


Figure 14.11 Naming the repository and choosing its type

Naming the repository is mandatory, followed by choosing a type: shared file system (fs), AWS S3, Azure's blob storage, and so on. Continuing with the example of choosing a local filesystem as our repository, click Shared File System to navigate to the next page. There, input the filesystem location and other requested properties (maximum and minimum bytes per second for snapshots and restores, chunk size, etc.). Click the Register button at the bottom of the page to create this repository.

14.4.3 Creating snapshots

Now that we've gone through the process of registering a snapshot repository, the next step is to create snapshots so the data is backed up to the repository we just created. There are a couple of ways to create a snapshot. Let's start with a simple manual technique: using the `_snapshot` API.

Listing 14.5 Creating a snapshot manually

```
PUT _snapshot/es_cluster_snapshot_repository/prod_snapshot_oct22
```

We ask the `_snapshot` API to create a snapshot named `prod_snapshot_oct22` under the repository `es_cluster_snapshot_repository`. This one-off manual snapshot backs up all the data (indexes, data streams, and cluster information) to the snapshot on the disk in the repository's filesystem.